

## 6.824 2017 Lecture 8: Zookeeper Case Study

Reading: "ZooKeeper: wait-free coordination for internet-scale systems", Patrick Hunt, Mahadev Konar, Flavio P. Junqueira, Benjamin Reed. Proceedings of the 2010 USENIX Annual Technical Conference.

Why are we reading this paper?

- Widely-used replicated state machine service

- Inspired by Chubby (Google's global lock service)

- Originally at Yahoo!, now outside too (Mesos, HBase, etc.)

- Open source

- As Apache project (<http://zookeeper.apache.org/>)

- Case study of building replicated services, given a Paxos/ZAB/Raft library

- Similar issues show up in lab 3

- API supports a wide-range of use cases

- Application that need a fault-tolerant "master" don't need to roll their own

- Zookeeper is generic enough that they should be able to use Zookeeper

- High performance

- Unlike lab 3's replicate key/value service

Motivation: many applications in datacenter cluster need to coordinate

- Example: GFS

- master has list of chunk servers for each chunk

- master decides which chunk server is primary

- etc.

- Other examples: YMB, Crawler, etc.

- YMB needs master to shard topics

- Crawler needs master that commands page fetching

- (e.g., a bit like the master in mapreduce)

- Applications also need to find each other

- MapReduce needs to know IP:PORT of GFS master

- Load balancer needs to know where web servers are

- Coordination service typically used for this purpose

Motivation: performance -- lab3

- dominated by Raft

- consider a 3-node Raft

- before returning to client, Raft performs

- leader persists log entry

- in parallel, leader send message to followers

- each follower persist log entry

- each follower responds

- > 2 disk writes and one round trip

- if magnetic disk:  $2 \times 10 \text{ msec} = 50 \text{ msg/sec}$

- if SSD:  $2 \times 2 \text{ msec} + 1 \text{ msec} = 200 \text{ msg/sec}$

- Zookeeper performs 21,000 msg/sec

- asynchronous calls

- allows pipelining

Alternative plan: develop fault-tolerant master for each application

- announce location via DNS

- OK, if writing master isn't complicated

- But, master often needs to be:

- fault tolerant

- every application figures how to use Raft?

- high performance

- every application figures how to make "read" operations fast?

- DNS propagation is slow

- fail-over will take a long time!

- Some application settle for single-point of failure

- E.g., GFS and MapReduce

- Less desirable

Zookeeper: a generic coordination service

- Design challenges:

- What API?

- How to make master fault tolerant?

- How to get good performance?

- Challenges interact

- good performance may influence API

- e.g., asynchronous interface to allow pipelining

Zookeeper API overview

- [diagram: ZooKeeper, client sessions, ZAB layer]

- replicated state machine

- several servers implementing the service

- operations are performed in global order
  - with some exceptions, if consistency isn't important
- the replicated objects are: znodes
  - hierarchy of znodes
    - named by pathnames
  - znodes contain \*metadata\* of application
    - configuration information
      - machines that participate in the application
        - which machine is the primary
    - timestamps
    - version number
- types of znodes:
  - regular
  - ephemeral
  - sequential: name + seqno
    - If n is the new znode and p is the parent znode, then the sequence value of n is never smaller than the value in the name of any other sequential znode ever created under p.

sessions

- clients sign into zookeeper
- session allows a client to fail-over to another Zookeeper service
  - client know the term and index of last completed operation (zxid)
  - send it on each request
    - service performs operation only if caught up with what client has seen
- sessions can timeout
  - client must refresh a session continuously
    - send a heartbeat to the server (like a lease)
  - ZooKeeper considers client "dead" if doesn't hear from a client
    - client may keep doing its thing (e.g., network partition)
      - but cannot perform other ZooKeeper ops in that session
- no analogue to this in Raft + Lab 3 KV store

Operations on znodes

- create(path, data, flags)
- delete(path, version)
  - if znode.version = version, then delete
- exists(path, watch)
- getData(path, watch)
- setData(path, data, version)
  - if znode.version = version, then update
- getChildren(path, watch)
- sync()
  - above operations are \*asynchronous\*
  - all operations are FIFO-ordered per client
  - sync waits until all preceding operations have been "propagated"

Check: can we just do this with lab 3's KV service?

- flawed plan: GFS master on startup does Put("gfs-master", my-ip:port)
  - other applications + GFS nodes do Get("gfs-master")
- problem: what if two master candidates' Put()'s race?
  - later Put() wins
  - each presumed master needs to read the key to see if it actually is the master
    - when are we assured that no delayed Put() thrashes us?
    - every other client must have seen our Put() -- hard to guarantee
- problem: when master fails, who decides to remove/update the KV store entry?
  - need some kind of timeout
  - so master must store tuple of (my-ip:port, timestamp)
    - and continuously Put() to refresh the timestamp
  - others poll the entry to see if the timestamp stops changing
- lots of polling + unclear race behavior -- complex
- ZooKeeper API has a better story: watches, sessions, atomic znode creation
  - + only one creation can succeed -- no Put() race
  - + sessions make timeouts easy -- no need to store and refresh explicit timestamps
  - + watches are lazy notifications -- avoids committing lots of polling reads

Ordering guarantees

- all write operations are totally ordered
  - if a write is performed by ZooKeeper, later writes from other clients see it
  - e.g., two clients create a znode, ZooKeeper performs them in some total order
- all operations are FIFO-ordered per client
- implications:
  - a read observes the result of an earlier write from the same client
  - a read observes some prefix of the writes, perhaps not including most recent write
    - > read can return stale data
  - if a read observes some prefix of writes, a later read observes that prefix too

Example "ready" znode:

- A failure happens
- A primary sends a stream of writes into Zookeeper
  - W1...Wn C(ready)
- The final write updates ready znode
  - > all preceding writes are visible
- The final write causes the watch to go off at backup
  - backup issues R(ready) R1...Rn
  - however, it will observe all writes because zookeeper will delay read until node has seen all txn that watch observed
- Lets say failure happens during R1 .. Rn, say after return Rj to client
  - primary deletes ready file -> watch goes off
  - watch alert is sent to client
  - client knows it must issue new R(ready) R1 ...Rn
- Nice property: high performance
  - pipeline writes and reads
  - can read from *\*any\** zookeeper node

Example usage 1: slow lock

```
acquire lock:
retry:
  r = create("app/lock", "", ephemeral)
  if r:
    return
  else:
    getData("app/lock", watch=True)

watch_event:
  goto retry

release lock: (voluntarily or session timeout)
delete("app/lock")
```

Example usage 2: "ticket" locks

```
acquire lock:
  n = create("app/lock/request-", "", ephemeral|sequential)
retry:
  requests = getChildren(l, false)
  if n is lowest znode in requests:
    return
  p = "request-%d" % n - 1
  if exists(p, watch = True)
    goto retry

watch_event:
  goto retry
```

Q: can watch\_event fire before lock it is the client's turn

A: yes

```
lock/request-10 <- current lock holder
lock/request-11 <- next one
lock/request-12 <- my request
```

if client associated with request-11 dies before it gets the lock, the watch event will fire but it isn't my turn yet.

Using locks

Not straight forward: a failure may cause your lock to be revoked

- client 1 acquires lock
  - starts doing its stuff
  - network partitions
  - zookeeper declares client 1 dead (but it isn't)
- client 2 acquires lock, but client 1 still believes it has it
  - can be avoided by setting timeouts correctly
  - need to disconnect client 1 session before ephemeral nodes go away
  - requires session heartbeats to be replicated to majority

N.B.: paper doesn't discuss this

For some cases, locks are a performance optimization

- for example, client 1 has a lock on crawling some urls
- client will do it 2 now, but that is fine

For other cases, locks are a building block

- for example, application uses it to build transaction
- the transactions are all-or-nothing
- we will see an example in the Frangipani paper

Zookeeper simplifies building applications but is not an end-to-end solution  
Plenty of hard problems left for application to deal with

Consider using Zookeeper in GFS

- I.e., replace master with Zookeeper

Application/GFS still needs all the other parts of GFS

- the primary/backup plan for chunks

- version numbers on chunks

- protocol for handling primary fail over

- etc.

With Zookeeper, at least master is fault tolerant

- And, won't run into split-brain problem

- Even though it has replicated servers

Implementation overview

- Similar to lab 3 (see last lecture)

- two layers:

- ZooKeeper services (K/V service)

- ZAB layer (Raft layer)

- Start() to insert ops in bottom layer

- Some time later ops pop out of bottom layer on each replica

- These ops are committed in the order they pop out

- on apply channel in lab 3

- the abdeliver() upcall in ZAB

Challenge: Duplicates client requests

- Scenario

- Primary receives client request, fails

- Client resends client request to new primary

- Lab 3:

- Table to detect duplicates

- Limitation: one outstanding op per client

- Problem problem: cannot pipeline client requests

- Zookeeper:

- Some ops are idempotent period

- Some ops are easy to make idempotent

- test-version-and-then-do-op

- e.g., include timestamp and version in setDataTXN

Challenge: Read operations

- Many operations are read operations

- they don't modify replicated state

- Must they go through ZAB/Raft or not?

- Can any replica execute the read op?

- Performance is slow if read ops go through Raft

Problem: read may return stale data if only master performs it

- The primary may not know that it isn't the primary anymore

- a network partition causes another node to become primary

- that partition may have processed write operations

- If the old primary serves read operations, it won't have seen those write ops

- => read returns stale data

Zookeeper solution: don't promise non-stale data (by default)

- Reads are allowed to return stale data

- Reads can be executed by any replica

- Read throughput increases as number of servers increases

- Read returns the last zxid it has seen

- So that new primary can catch up to zxid before serving the read

- Avoids reading from past

- Only sync-read() guarantees data is not stale

Sync optimization: avoid ZAB layer for sync-read

- must ensure that read observes last committed txn

- leader puts sync in queue between it and replica

- if ops ahead of in the queue commit, then leader must be leader

- otherwise, issue null transaction

- in same spirit read optimization in Raft paper

- see last par section 8 of raft paper

Performance (see table 1)

- Reads inexpensive

- Q: Why more reads as servers increase?

- Writes expensive

- Q: Why slower with increasing number of servers?

- Quick failure recovery (figure 8)

- Decent throughput even while failures happen

References:

- ZAB: <http://dl.acm.org/citation.cfm?id=2056409>

<https://zookeeper.apache.org/>

<https://cs.brown.edu/~mph/Herlihy91/p124-herlihy.pdf> (wait free, universal objects, etc.)