

6.824 2018 Lecture 12: Spark Case Study

Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing
Zaharia, Chowdhury, Das, Dave, Ma, McCauley, Franklin, Shenker, Stoica
NSDI 2012

Today: more distributed computations

Case study: Spark

Why are we reading Spark?

- Widely-used for datacenter computations
- popular open-source project, hot startup (Databricks)
- Support iterative applications better than MapReduce
- Interesting fault tolerance story
- ACM doctoral thesis award

MapReduce make life of programmers easy

It handles:

- Communication between nodes
- Distribute code
- Schedule work
- Handle failures

But restricted programming model

Some apps don't fit well with MapReduce

Many algorithms are iterative

Page-rank is the classic example

- compute a rank for each document, based on how many docs point to it
- the rank is used to determine the place of the doc in the search results

on each iteration, each document:

- sends a contribution of r/n to its neighbors
 - where r is its rank and n is its number of neighbors.
- update rank to $\alpha/N + (1 - \alpha) * \text{Sum}(c_i)$,
 - where the sum is over the contributions it received
 - N is the total number of documents.

big computation:

- runs over all web pages in the world
- even with many machines it takes long time

MapReduce and iterative algorithms

MapReduce would be good at one iteration of the algorithms

Each map does part of the documents

Reduce for update the rank of a particular doc

But what to do for the next iteration?

Write results to storage

Start a new MapReduce job for the next iteration

Expensive

But fault tolerant

Challenges

Better programming model for iterative computations

Good fault tolerance story

One solution: use DSM

Good for iterative programming

ranks can be in shared memory

workers can update and read

Bad for fault tolerance

typical plan: checkpoint state of memory

make a checkpoint every hour of memory

expensive in two ways:

write shared memory to storage during computation

redo all work since last checkpoint after failure

Spark more MapReduce flavor

Restricted programming model, but more powerful than MapReduce

Good fault tolerance plan

Better solution: keep data in memory

Pregel, Dryad, Spark, etc.

In Spark

Data is stored in data sets (RDDs)

"Persist" the RDD in memory

Next iteration can refer to the RDD

Other opportunities

Interactive data exploration

Run queries over the persisted RDDs

Like to have something SQL-like
A join operator over RDDs

Core idea in Spark: RDDs

RDDs are immutable --- you cannot update them

RDDs support transformations and actions

Transformations: compute a new RDD from existing RDDs

map, reduceByKey, filter, join, ..

transformations are lazy: don't compute result immediately

just a description of the computation

Actions: for when results are needed

counts result, collect results, get a specific value

Example use:

```
lines = spark.textFile("hdfs://...")
```

```
errors = lines.filter(_.startsWith("ERROR"))    // lazy!
```

```
errors.persist()    // no work yet
```

```
errors.count()    // an action that computes a result
```

```
// now errors is materialized in memory
```

```
// partitioned across many nodes
```

```
// Spark, will try to keep in RAM (will spill to disk when RAM is full)
```

Reuse of an RDD

```
errors.filter(_.contains("MySQL")).count()
```

```
// this will be fast because reuses results computed by previous fragment
```

```
// Spark will schedule jobs across machines that hold partition of errors
```

Another reuse of RDD

```
errors.filter(_.contains("HDFS")).map(_.split('\t')(3)).collect()
```

RDD lineage

Spark creates a lineage graph on an action

Graphs describe the computation using transformations

lines -> filter w ERROR -> errors -> filter w. HDFS -> map -> timed fields

Spark uses the lineage to schedule job

Transformation on the same partition form a stage

Joins, for example, are a stage boundary

Need to reshuffle data

A job runs a single stage

pipeline transformation within a stage

Schedule job where the RDD partition is

Lineage and fault tolerance

Great opportunity for *efficient* fault tolerance

Let's say one machine fails

Want to recompute *only* its state

The lineage tells us what to recompute

Follow the lineage to identify all partitions needed

Recompute them

For last example, identify partitions of lines missing

Trace back from child to parents in lineage

All dependencies are "narrow"

each partition is dependent on one parent partition

Need to read the missing partition of lines

recompute the transformations

RDD implementation

list of partitions

list of (parent RDD, wide/narrow dependency)

narrow: depends on one parent partition (e.g., map)

wide: depends on several parent partitions (e.g., join)

function to compute (e.g., map, join)

partitioning scheme (e.g., for file by block)

computation placement hint

Each transformation takes (one or more) RDDs, and outputs the transformed RDD.

Q: Why does an RDD carry metadata on its partitioning? A: so transformations that depend on multiple RDDs know whether they need to shuffle data (wide dependency) or not (narrow). Allows users control over locality and reduces shuffles.

Q: Why the distinction between narrow and wide dependencies? A: In case of failure. Narrow dependency only depends on a few partitions that need to be recomputed. Wide dependency might require an entire RDD

Example: PageRank (from paper):

```
// Load graph as an RDD of (URL, outlinks) pairs
val links = spark.textFile(...).map(...).persist() // (URL, outlinks)
var ranks = // RDD of (URL, rank) pairs
for (i <- 1 to ITERATIONS) {
  // Build an RDD of (targetURL, float) pairs
  // with the contributions sent by each page
  val contribs = links.join(ranks).flatMap {
    (url, (links, rank)) => links.map(dest => (dest, rank/links.size))
  }
  // Sum contributions by URL and get new ranks
  ranks = contribs.reduceByKey((x,y) => x+y)
    .mapValues(sum => a/N + (1-a)*sum)
}
```

Lineage for PageRank

See figure 3

Each iteration creates two new RDDs:

ranks0, ranks1, etc.

contribs0, contribs1, etc.

Long lineage graph!

Risky for fault tolerance.

One node fails, much recomputation

Solution: user can replicate RDD

Programmer pass "reliable" flag to persist()

e.g., call ranks.persist(RELIABLE) every N iterations

Replicates RDD in memory

With REPLICATE flag, will write to stable storage (HDFS)

Impact on performance

if user frequently persist w/REPLICATE, fast recovery, but slower execution

if infrequently, fast execution but slow recovery

Q: Is persist a transformation or an action? A: neither. It doesn't create a new RDD, and doesn't cause materialization. It's an instruction to the scheduler.

Q: By calling persist without flags, is it guaranteed that in case of fault that RDD wouldn't have to be recomputed? A: No. There is no replication, so a node holding a partition could fail. Replication (either in RAM or in stable storage) is necessary

Currently only manual checkpointing via calls to persist. Q: Why implement checkpointing? (it's expensive) A: Long lineage could cause large recovery time. Or when there are wide dependencies a single failure might require many partition re-computations.

Q: Can Spark handle network partitions? A: Nodes that cannot communicate with scheduler will appear dead. The part of the network that can be reached from scheduler can continue computation, as long as it has enough data to start the lineage from (if all replicas of a required partition cannot be reached, cluster cannot make progress)

What happens when there isn't enough memory?

LRU (Least Recently Used) on partitions

first on non-persisted

then persisted (but they will be available on disk. makes sure user cannot overbook RAM)

User can have control on order of eviction via "persistence priority"

No reason not to discard non-persisted partitions (if they've already been used)

Performance

Degrades to "almost" MapReduce behavior

In figure 7, logistic regression on 100 Hadoop nodes takes 76-80 seconds

In figure 12, logistic regression on 25 Spark nodes (with no partitions allowed in memory) takes 68.8

Difference could be because MapReduce uses replicated storage after reduce, but Spark by default only spills to local disk

no network latency and I/O load on replicas.

no architectural reason why MR would be slower than Spark for non-iterative work

or for iterative work that needs to go to disk

no architectural reason why Spark would ever be slower than MR

Discussion

Spark targets batch, iterative applications

Spark can express other models

MapReduce, Pregel

Cannot incorporate new data as it comes in

But see Streaming Spark

Spark not good for building key/value store

Like MapReduce, and others
RDDs are immutable

References

<http://spark.apache.org/>

<http://www.cs.princeton.edu/~chazelle/courses/BIB/pagerank.htm>