

6.824 2015 Lecture 16: Scaling Memcache at Facebook

Scaling Memcache at Facebook, by Nishtala et al, NSDI 2013

why are we reading this paper?

it's an experience paper, not about new ideas/techniques

three ways to read it:

cautionary tale of problems from not taking consistency seriously

impressive story of super high capacity from mostly-off-the-shelf s/w

fundamental struggle between performance and consistency

we can argue with their design, but not their success

how do web sites scale up with growing load?

a typical story of evolution over time:

1. one machine, web server, application, DB

DB stores on disk, crash recovery, transactions, SQL

application queries DB, formats, HTML, &c

but the load grows, your PHP application takes too much CPU time

2. many web FEs, one shared DB

an easy change, since web server + app already separate from storage

FEs are stateless, all sharing (and concurrency control) via DB

but the load grows; add more FEs; soon single DB server is bottleneck

3. many web FEs, data sharded over cluster of DBs

partition data by key over the DBs

app looks at key (e.g. user), chooses the right DB

good DB parallelism if no data is super-popular

painful -- cross-shard transactions and queries probably don't work

hard to partition too finely

but DBs are slow, even for reads, why not cache read requests?

4. many web FEs, many caches for reads, many DBs for writes

cost-effective b/c read-heavy and memcached 10x faster than a DB

memcached just an in-memory hash table, very simple

complex b/c DB and memcacheds can get out of sync

(next bottleneck will be DB writes -- hard to solve)

the big facebook infrastructure picture

lots of users, friend lists, status, posts, likes, photos

fresh/consistent data apparently not critical

because humans are tolerant?

high load: billions of operations per second

that's 10,000x the throughput of one DB server

multiple data centers (at least west and east coast)

each data center -- "region":

"real" data sharded over MySQL DBs

memcached layer (mc)

web servers (clients of memcached)

each data center's DBs contain full replica

west coast is master, others are slaves via MySQL async log replication

how do FB apps use mc?

read:

v = get(k) (computes hash(k) to choose mc server)

if v is nil {

v = fetch from DB

set(k, v)

}

write:

v = new value

send k,v to DB

delete(k)

application determines relationship of mc to DB

mc doesn't know anything about DB

FB uses mc as a "look-aside" cache

real data is in the DB

cached value (if any) should be same as DB

what does FB store in mc?

paper does not say

maybe userID -> name; userID -> friend list; postID -> text; URL -> likes

basically copies of data from DB

paper lessons:

look-aside is much trickier than it looks -- consistency

paper is trying to integrate mutually-oblivious storage layers

cache is critical:

not really about reducing user-visible delay

- mostly about surviving huge load!
- cache misses and failures can create intolerable DB load
- they can tolerate modest staleness: no freshness guarantee
- stale data nevertheless a big headache
 - want to avoid unbounded staleness (e.g. missing a delete() entirely)
 - want read-your-own-writes
 - each performance fix brings a new source of staleness
- huge "fan-out" => parallel fetch, in-cast congestion

let's talk about performance first

- majority of paper is about avoiding stale data
- but staleness only arose from performance design

performance comes from parallel get()s by many mc servers

- driven by parallel processing of HTTP requests by many web servers
- two basic parallel strategies for storage: partition vs replication

will partition or replication yield most mc throughput?

- partition: server i, key k -> mc server hash(k)
- replicate: server i, key k -> mc server hash(i)
- partition is more memory efficient (one copy of each k/v)
- partition works well if no key is very popular
- partition forces each web server to talk to many mc servers (overhead)
- replication works better if a few keys are very popular

performance and regions (Section 5)

Q: what is the point of regions -- multiple complete replicas?

- lower RTT to users (east coast, west coast)
- parallel reads of popular data due to replication
- (note DB replicas help only read performance, no write performance)
- maybe hot replica for main site failure?

Q: why not partition users over regions?

- i.e. why not east-coast users' data in east-coast region, &c
- social net -> not much locality
- very different from e.g. e-mail

Q: why OK performance despite all writes forced to go to the master region?

- writes would need to be sent to all regions anyway -- replicas
- users probably wait for round-trip to update DB in master region
 - only 100ms, not so bad
- users do not wait for all effects of writes to finish
 - i.e. for all stale cached values to be deleted

performance within a region (Section 4)

multiple mc clusters *within* each region

- cluster == complete set of mc cache servers
 - i.e. a replica, at least of cached data

why multiple clusters per region?

- why not add more and more mc servers to a single cluster?
- 1. adding mc servers to cluster doesn't help single popular keys
 - replicating (one copy per cluster) does help
- 2. more mcs in cluster -> each client req talks to more servers
 - and more in-cast congestion at requesting web servers
 - client requests fetch 20 to 500 keys! over many mc servers
 - MUST request them in parallel (otherwise total latency too large)
 - so all replies come back at the same time
 - network switches, NIC run out of buffers
- 3. hard to build network for single big cluster
 - uniform client/server access
 - so cross-section b/w must be large -- expensive
 - two clusters -> 1/2 the cross-section b/w

but -- replicating is a waste of RAM for less-popular items

- "regional pool" shared by all clusters
- unpopular objects (no need for many copies)
- decided by *type* of object
- frees RAM to replicate more popular objects

bringing up new mc cluster was a serious performance problem

- new cluster has 0% hit rate
- if clients use it, will generate big spike in DB load
 - if ordinarily 1% miss rate, and (let's say) 2 clusters,
 - adding "cold" third cluster will causes misses for 33% of ops.

- i.e. 30x spike in DB load!
- thus the clients of new cluster first `get()` from existing cluster (4.3)
- and `set()` into new cluster
- basically lazy copy of existing cluster to new cluster
- better 2x load on existing cluster than 30x load on DB

important practical networking problems:

- n^2 TCP connections is too much state
- thus UDP for client `get()`s
- UDP is not reliable or ordered
- thus TCP for client `set()`s
- and mcrouter to reduce n in n^2
- small request per packet is not efficient (for TCP or UDP)
- per-packet overhead (interrupt &c) is too high
- thus mcrouter batches many requests into each packet

mc server failure?

- can't have DB servers handle the misses -- too much load
- can't shift load to one other mc server -- too much
- can't re-partition all data -- time consuming
- Gutter -- pool of idle servers, clients only use after mc server fails

The Question:

- why don't clients send invalidates to Gutter servers?
- my guess: would double `delete()` traffic
- and send too many `delete()`s to small gutter pool
- since any key might be in the gutter pool

thundering herd

- one client updates DB and `delete()`s a key
- lots of clients `get()` but miss
- they all fetch from DB
- they all `set()`
- not good: needless DB load
- mc gives just the first missing client a "lease"
- lease = permission to refresh from DB
- mc tells others "try `get()` again in a few milliseconds"
- effect: only one client reads the DB and does `set()`
- others re-try `get()` later and hopefully hit

let's talk about consistency now

the big truth

- hard to get both consistency (== freshness) and performance
- performance for reads = many copies
- many copies = hard to keep them equal

what is their consistency goal?

- *not* read sees latest write
- since not guaranteed across clusters
- more like "not more than a few seconds stale"
- i.e. eventual
- *and* writers see their own writes
- read-your-own-writes is a big driving force

first, how are DB replicas kept consistent across regions?

- one region is master
- master DBs distribute log of updates to DBs in slave regions
- slave DBs apply
- slave DBs are complete replicas (not caches)
- DB replication delay can be considerable (many seconds)

how do we feel about the consistency of the DB replication scheme?

- good: eventual consistency, b/c single ordered write stream
- bad: longish replication delay -> stale reads

how do they keep mc content consistent w/ DB content?

1. DBs send invalidates (`delete()`s) to all mc servers that might cache
2. writing client also invalidates mc in local cluster

for read-your-writes

why did they have consistency problems in mc?

- client code to copy DB to mc wasn't atomic:
- 1. writes: DB update ... mc `delete()`
- 2. read miss: DB read ... mc `set()`

so *concurrent* clients had races

what were the races and fixes?

Race 1:

- k not in cache
- C1 get(k), misses
- C1 v = read k from DB
- C2 updates k in DB
- C2 and DB delete(k) -- does nothing
- C1 set(k, v)
- now mc has stale data, delete(k) has already happened
- will stay stale indefinitely, until key is next written
- solved with leases -- C1 gets a lease, but C2's delete() invalidates lease, so mc ignores C1's set
- key still missing, so next reader will refresh it from DB

Race 2:

- during cold cluster warm-up
- remember clients try get() in warm cluster, copy to cold cluster
- k starts with value v1
- C1 updates k to v2 in DB
- C1 delete(k) -- in cold cluster
- C2 get(k), miss -- in cold cluster
- C2 v1 = get(k) from warm cluster, hits
- C2 set(k, v1) into cold cluster
- now mc has stale v1, but delete() has already happened
- will stay stale indefinitely, until key is next written
- solved with two-second hold-off, just used on cold clusters
- after C1 delete(), cold ignores set()s for two seconds
- by then, delete() will propagate via DB to warm cluster

Race 3:

- k starts with value v1
- C1 is in a slave region
- C1 updates k=v2 in master DB
- C1 delete(k) -- local region
- C1 get(k), miss
- C1 read local DB -- sees v1, not v2!
- later, v2 arrives from master DB
- solved by "remote mark"
- C1 delete() marks key "remote"
- get()/miss yields "remote"
- tells C1 to read from *master* region
- "remote" cleared when new data arrives from master region

Q: aren't all these problems caused by clients copying DB data to mc?
why not instead have DB send new values to mc, so clients only read mc?
then there would be no racing client updates &c, just ordered writes

A:

1. DB doesn't generally know how to compute values for mc
generally client app code computes them from DB results,
i.e. mc content is often not simply a literal DB record
2. would increase read-your-own writes delay
3. DB doesn't know what's cached, would end up sending lots
of values for keys that aren't cached

PNUTS does take this alternate approach of master-updates-all-copies

FB/mc lessons for storage system designers?

- cache is vital to throughput survival, not just a latency tweak
- need flexible tools for controlling partition vs replication
- need better ideas for integrating storage layers with consistency