

Dynamo: Amazon's Highly Available Key-value Store
DeCandia et al, SOSP 2007

Why are we reading this paper?

- Database, eventually consistent, write any replica
- Like Bayou, with reconciliation
- Like Parameter Server, but geo-distributed
- A surprising design.

- A real system: used for e.g. shopping cart at Amazon
- More available than PNUTS, Spanner, FB MySQL, &c
- Less consistent than PNUTS, Spanner, FB MySQL &c
- Influential design; inspired e.g. Cassandra
- 2007: before PNUTS, before Spanner

Their Obsessions

- SLA, e.g. 99.9th percentile of delay < 300 ms
- constant failures
- "data centers being destroyed by tornadoes"
- "always writeable"

Big picture

- [lots of data centers, Dynamo nodes]
- each item replicated at a few random nodes, by key hash

Why replicas at just a few sites? Why not replica at every site?

- with two data centers, site failure takes down 1/2 of nodes
 - so need to be careful that *everything* replicated at *both* sites
- with 10 data centers, site failure affects small fraction of nodes
 - so just need copies at a few sites

Where to place data -- consistent hashing

- [ring, and physical view of servers]
- node ID = random
- key ID = hash(key)
- coordinator: successor of key
 - clients send puts/gets to coordinator
- replicas at successors -- "preference list"
- coordinator forwards puts (and gets...) to nodes on preference list

Consequences of mostly remote access (since no guaranteed local copy)

- most puts/gets may involve WAN traffic -- high delays
- the quorums will cut the tail end --- see below
- but can survive data centers going down

Why consistent hashing?

- Pro
 - naturally somewhat balanced
 - decentralized -- both lookup and join/leave
- Con (section 6.2)
 - not really balanced (why not?), need virtual nodes
 - hard to control placement (balancing popular keys, spread over sites)
 - join/leave changes partition, requires data to shift

Failures

- Tension: temporary or permanent failure?
 - node unreachable -- what to do?
 - if temporary, store new puts elsewhere until node is available
 - if permanent, need to make new replica of all content
- Dynamo itself treats all failures as temporary

Consequences of "always writeable"

- always writeable => no master! must be able to write locally.
 - idea 1: sloppy quorums
- always writeable + failures = conflicting versions
 - idea 2: eventual consistency
 - idea 1 avoids inconsistencies when there are no failures

Idea #1: sloppy quorum

- try to get consistency benefits of single master if no failures
 - but allows progress even if coordinator fails, which PNUTS does not
- when no failures, send reads/writes through single node
 - the coordinator
 - causes reads to see writes in the usual case

but don't insist! allow reads/writes to any replica if failures

Temporary failure handling: quorum

goal: do not block waiting for unreachable nodes

goal: put should always succeed

goal: get should have high prob of seeing most recent put(s)

quorum: $R + W > N$

never wait for all N

but R and W will overlap

cuts tail off delay distribution and tolerates some failures

N is first N *reachable* nodes in preference list

each node pings successors to keep rough estimate of up/down

"sloppy" quorum, since nodes may disagree on reachable

sloppy quorum means R/W overlap *not guaranteed*

coordinator handling of put/get:

sends put/get to first N reachable nodes, in parallel

put: waits for W replies

get: waits for R replies

if failures aren't too crazy, get will see all recent put versions

When might this quorum scheme *not* provide R/W intersection?

What if a put() leaves data far down the ring?

after failures repaired, new data is beyond N?

that server remembers a "hint" about where data really belongs

forwards once real home is reachable

also -- periodic "merkle tree" sync of key range

Idea #2: eventual consistency

accept writes at any replica

allow divergent replicas

allow reads to see stale or conflicting data

resolve multiple versions when failures go away

latest version if no conflicting updates

if conflicts, reader must merge and then write

like Bayou and Ficus -- but in a DB

Unhappy consequences of eventual consistency

May be no unique "latest version"

Read can yield multiple conflicting versions

Application must merge and resolve conflicts

No atomic operations (e.g. no PNUTS test-and-set-write)

How can multiple versions arise?

Maybe a node missed the latest write due to network problem

So it has old data, should be superseded by newer put()s

get() consults R, will likely see newer version as well as old

How can *conflicting* versions arise?

N=3 R=2 W=2

shopping cart, starts out empty ""

preference list n1, n2, n3, n4

client 1 wants to add item X

get() from n1, n2, yields ""

n1 and n2 fail

put("X") goes to n3, n4

n1, n2 revive

client 3 wants to add Y

get() from n1, n2 yields ""

put("Y") to n1, n2

client 3 wants to display cart

get() from n1, n3 yields two values!

"X" and "Y"

neither supersedes the other -- the put()s conflicted

How should clients resolve conflicts on read?

Depends on the application

Shopping basket: merge by taking union?

Would un-delete items removed

Weaker than Bayou (which gets deletion right), but simpler

Some apps probably can use latest wall-clock time

E.g. if I'm updating my password

Simpler for apps than merging

Write the merged result back to Dynamo

Programming API

All objects are immutable
- get(k) may return multiple versions, along with "context"
- put(k, v, context)
creates a new version of k, attaching context
The context is used to merge and keep track of dependencies, and
detect how conflicts. It consists of a VV of the object.

Version vectors

Example tree of versions:

```
[a:1]
|
+-----|
      [a:1,b:2]
```

VVs indicate v2 supersedes v1

Dynamo nodes automatically drop [a:1] in favor of [a:1,b:2]

Example:

```
[a:1]
|
+-----|
|      [a:1,b:2]
|
[a:2]
```

Client must merge

Won't the VVs get big?

Yes, but slowly, since key mostly served from same N nodes

Dynamo deletes least-recently-updated entry if VV has > 10 elements

Impact of deleting a VV entry?

won't realize one version subsumes another, will merge when not needed:

put@b: [b:4]

put@a: [a:3, b:4]

forget b:4: [a:3]

now, if you sync w/ [b:4], looks like a merge is required

forgetting the oldest is clever

since that's the element most likely to be present in other branches

so if it's missing, forces a merge

forgetting *newest* would erase evidence of recent difference

Is client merge of conflicting versions always possible?

Suppose we're keeping a counter, x

x starts out 0

incremented twice

but failures prevent clients from seeing each others' writes

After heal, client sees two versions, both x=1

What's the correct merge result?

Can the client figure it out?

What if two clients concurrently write w/o failure?

e.g. two clients add diff items to same cart at same time

Each does get-modify-put

They both see the same initial version

And they both send put() to same coordinator

Will coordinator create two versions with conflicting VVs?

We want that outcome, otherwise one was thrown away

Paper doesn't say, but coordinator could detect problem via put() context

Permanent server failures / additions?

Admin manually modifies the list of servers

System shuffles data around -- this takes a long time!

The Question:

It takes a while for notice of added/deleted server to become known
to all other servers. Does this cause trouble?

Deleted server might get put()s meant for its replacement.

Deleted server might receive get()s after missing some put()s.

Added server might miss some put()s b/c not known to coordinator.

Added server might serve get()s before fully initialized.

Dynamo probably will do the right thing:

Quorum likely causes get() to see fresh data as well as stale.

Replica sync (4.7) will fix missed get()s.

Is the design inherently low delay?

No: client may be forced to contact distant coordinator

No: some of the R/W nodes may be distant, coordinator must wait

What parts of design are likely to help limit 99.9th pctl delay?

This is a question about variance, not mean
Bad news: waiting for multiple servers takes *max* of delays, not e.g. avg
Good news: Dynamo only waits for W or R out of N
cuts off tail of delay distribution
e.g. if nodes have 1% chance of being busy with something else
or if a few nodes are broken, network overloaded, &c

No real Eval section, only Experience

How does Amazon use Dynamo?
shopping cart (merge)
session info (maybe Recently Visited &c?) (most recent TS)
product list (mostly r/o, replication for high read throughput)

They claim main advantage of Dynamo is flexible N, R, W
What do you get by varying them?
N-R-W
3-2-2 : default, reasonable fast R/W, reasonable durability
3-3-1 : fast W, slow R, not very durable, not useful?
3-1-3 : fast R, slow W, durable
3-3-3 : ??? reduce chance of R missing W?
3-1-1 : not useful?

They had to fiddle with the partitioning / placement / load balance (6.2)
Old scheme:
Random choice of node ID meant new node had to split old nodes' ranges
Which required expensive scans of on-disk DBs
New scheme:
Pre-determined set of Q evenly divided ranges
Each node is coordinator for a few of them
New node takes over a few entire ranges
Store each range in a file, can xfer whole file

How useful is ability to have multiple versions? (6.3)
I.e. how useful is eventual consistency
This is a Big Question for them
6.3 claims 0.001% of reads see divergent versions
I believe they mean conflicting versions (not benign multiple versions)
Is that a lot, or a little?
So perhaps 0.001% of writes benefitted from always-writeable?
I.e. would have blocked in primary/backup scheme?
Very hard to guess:
They hint that the problem was concurrent writers, for which
better solution is single master
But also maybe their measurement doesn't count situations where
availability would have been worse if single master

Performance / throughput (Figure 4, 6.1)
Figure 4 says average 10ms read, 20 ms writes
the 20 ms must include a disk write
10 ms probably includes waiting for R/W of N
Figure 4 says 99.9th pctl is about 100 or 200 ms
Why?
"request load, object sizes, locality patterns"
does this mean sometimes they had to wait for coast-coast msg?

Puzzle: why are the average delays in Figure 4 and Table 2 so low?
Implies they rarely wait for WAN delays
But Section 6 says "multiple datacenters"
You'd expect *most* coordinators and most nodes to be remote!
Maybe all datacenters are near Seattle?
Maybe because coordinators can be any node in the preference list?
See last paragraph of 5
Maybe W-1 copies in N are close by?

Wrap-up
Big ideas:
eventual consistency
always writeable despite failures
allow conflicting writes, client merges
Awkward model for some applications (stale reads, merges)
this is hard for us to tell from paper
Maybe a good way to get high availability + no blocking on WAN
Parameter Server uses similar ideas for ML applications
no single master, conflicting writes okay
No agreement on whether eventual consistency is good for storage systems